


Lesson objectives

- Learn what SparkR is
- Learn why you would use SparkR
- List the features of SparkR
- Understand the interfaces into SparkR

What is SparkR?

- An R package to use Apache Spark from R



- SparkR implements distributed dataframes
- SparkR supports operations like selection, filtering, aggregation etc.
- SparkR also supports MLIB

Why use SparkR?

- R has usability issues with big data workflows
- R is very resource constrained so limited optimizations
- R has restricted machine learning

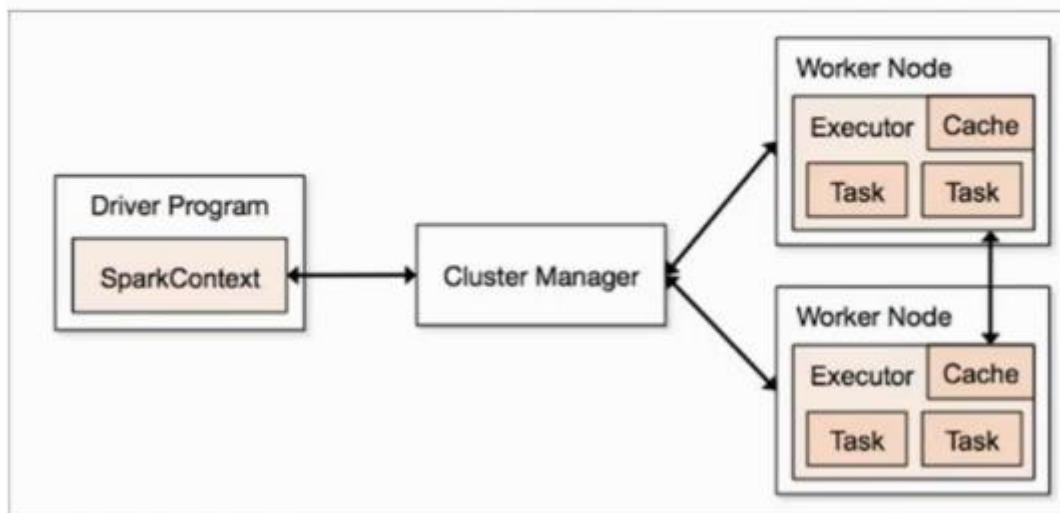
Features of SparkR

- Scalability to many cores and machines
- Data Frame Optimizations
- Data Sources API
- RDDs as Distributed Lists
- Serializing closures
- Using existing R packages

Interfaces

- Spark shell
- SparkR shell
- [Rstudio](#)
- Notebooks
- Data Scientist Workbench (an interactive service providing notebooks interface)

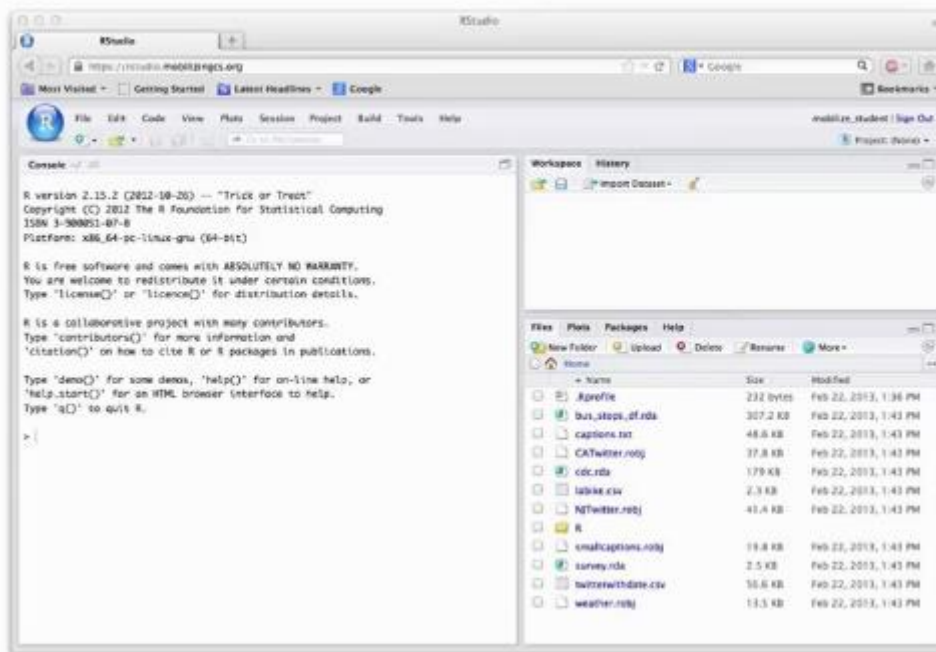
SparkContext



Spark SQL

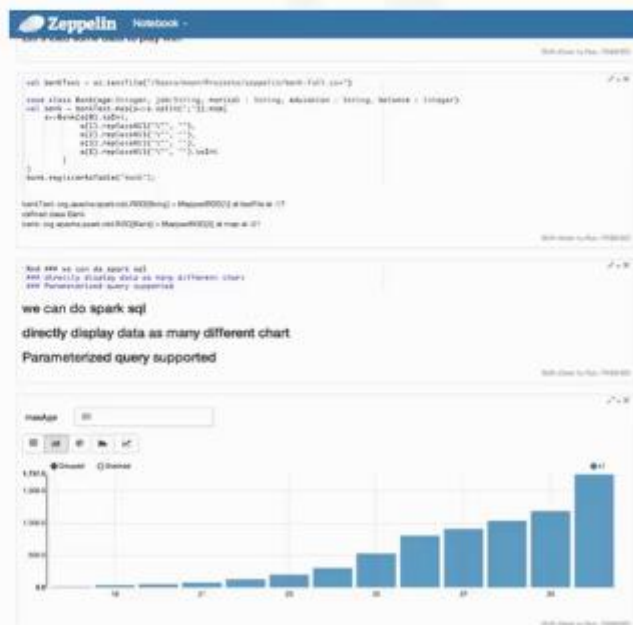
- Spark SQL is a Spark module for structured data processing
- The entry point into all relational functionality in Spark is the [SQLContext](#)
- There are several ways to interact with Spark SQL including SQL, the DataFrames API and the Datasets API:
 - SQL
 - JDBC/ODBC
 - Datasets with a strongly-typed LINQ-like Query DSL
- SparkR works solely with DataFrames which requires Spark SQL

RStudio

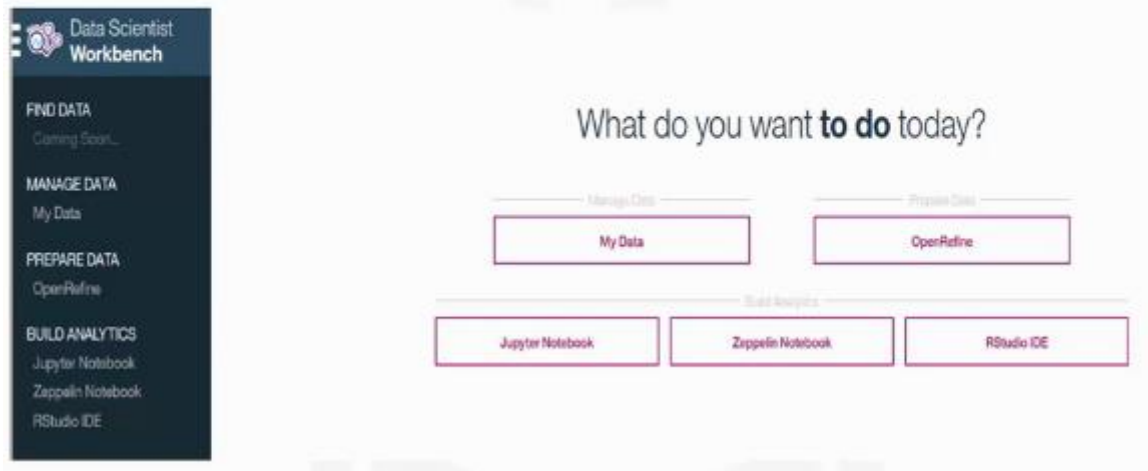


Notebooks

- Zeppelin notebook



Data Scientist Workbench (DSWB)



Lesson objectives

- Understand how to use dataframes
- Learn to select data
- Learn to filter data
- Learn to aggregate data
- Learn to operate on columns
- Understand how to write SQL queries

Dataframes

- Why dataframes
- Creating dataframes
 - Local dataframes
 - Dataframes from raw data
 - Dataframes from data sources

Name Conflicts

- It is possible to have a name conflict between functions in SparkR and R
- The following have a conflict:
 - cov in package:stats
 - filter in package:stats
 - sample in package:base
 - table in package:base
- It is also possible to have name conflicts between dplyr and SparkR

Select columns

- `select`

```
SparkR::head(select(cars, cars$mpg))
```

```
Out[3]:      mpg
1  21.0
2  21.0
3  22.8
4  21.4
5  18.7
6  18.1
```

- `selectExpr`

```
SparkR::head(selectExpr(df, "Day", "Day + 99", "Day * 3"))
```

	Day	(Day + 99)	(Day * 3)
1	1	100	3
2	2	101	6
3	3	102	9
4	4	103	12
5	5	104	15
6	6	105	18

RED DATA

Operating on columns

- Operating on a specific column using `$`

```
cars$mpg <- cars$mpg/3.78541178
```

```
Out[4]:      mpg  cyl disp  hp drat   wt  qsec vs  am gear carb
1  5.547613    6  160  110  3.90  2.620  16.46  0   1    4    4
2  5.547613    6  160  110  3.90  2.875  17.02  0   1    4    4
3  6.023123    4  108   93  3.85  2.320  18.61  1   1    4    1
4  5.653282    6  258  110  3.08  3.215  19.44  1   0    3    1
5  4.940017    8  360  175  3.15  3.440  17.02  0   0    3    2
6  4.781514    6  225  105  2.76  3.460  20.22  1   0    3    1
```

- Operating on multiple columns using `[]`

```
[, c("you", "me")]
```

Filter (select rows)

- filter by condition for one column

```
filtered <- SparkR::head(SparkR::filter(cars, df$cyl == 6)  
collect(filtered)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
4	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1

- filter by condition for multiple columns

- These two expression are same:

```
SparkR::head(SparkR::filter("cars, df$cyl == 6 and cars, cars$mpg < 20")
```

```
SparkR::head(SparkR::filter(cars, df$cyl == 6). SparkR::filter(cars, cars$mpg  
< 20)
```

SparkSQL (row operations)

- Dataframes as SQL tables

```
registerTempTable(cars, "cars")
```

```
SparkR::head(sql(sqlContext, "SELECT * FROM cars WHERE cyl > 6"))
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	18.7	8	360.0	175	3.15	3.44	17.02	0	0	3	2
2	14.3	8	360.0	245	3.21	3.57	15.84	0	0	3	4
3	16.4	8	275.8	180	3.07	4.07	17.40	0	0	3	3

Aggregate data

- groupBy for one column (count, sum, average)

- groupBy(cars, cars\$mpg), count = n(cars\$mpg))
- groupBy(cars, cars\$mpg), sum = sum(cars\$mpg))
- groupBy(cars, cars\$mpg), average = avg(cars\$hp))

- groupBy for multiple columns (count, sum, average)

- groupBy(cars, "'class', 'year'"), avg = n('hwy'))

Lesson objectives

- Understand machine learning
- Learn how to use GLM model

Machine learning

- Example machine learning workflow



- Observation = rows in a dataframe
- Feature = columns in a dataframe

MLlib and ML

- MLlib is Spark's machine learning (ML) library
- It divides into two packages:
 - spark.mllib contains the original API built on top of RDDs.
 - spark.ml provides higher-level API built on top of DataFrames
- Concepts in pipelines:
 - DataFrame
 - Transformer
 - Estimator
 - Pipeline
 - Parameter

Pipelines components

- Transformers
 - A Transformer is an abstraction that includes feature transformers and learned models
- Estimators
 - An Estimator abstracts the concept of a learning algorithm that trains on data
- Pipeline
 - a sequence of algorithms to process and learn from data

Implementing Linear Models in SparkR

- Prepare and load data
 - Load the data
 - Read the data in a Spark dataframe
 - Create factors
- Train the model
 - Formula
 - Dataset
 - Model
- Evaluate the model
 - Baseline reference
 - Predict()
- Implement model
 - Iterate through the training dataset
 - Find acceptable model

GLM data load

- Download and save the data files using R
 - `population_data_files_url <- 'http://www2.census.gov/acs2013_1yr/pums/csv_pus.zip'`
 - `library(RCurl)`
 - `population_data_file <- getBinaryURL(population_data_files_url)`
 - `population_data_file_path <- '/nfs/data/2013-acs/csv_pus.zip'`
 - `population_data_file_local <- file(population_data_file_path, open = "wb")`
- `sparkContext`
 - `sc <- sparkR.init()`
- `sqlContext`
 - `sqlContext <- sparkRSQL.init(sc)`

GLM data prep

- Convert any categorical variable from a numeric variable into a factor for example
 - `housing_df$ST <- cast(housing_df$ST, "string")`
- Sanitize the data for example remove nulls
 - `housing_with_valp_df <- filter(housing_df, isNotNull(housing_df$VALP) & isNotNull(housing_df$TAXP) & isNotNull(housing_df$INSP) & isNotNull(housing_df$ACR))`
- Create training and test data
 - `housing_df_test <- sample(housing_with_valp_df, FALSE, 0.1)`

GLM model training

- `glm()`
 - Fit a gaussian GLM model over the dataset.
 - `model <- glm(Sepal_Length ~ Sepal_Width + Species, data = df, family = "gaussian")`
- `summary()`
 - Model summary are returned in a similar format to R's native `glm()`.
 - `summary(model)`
 - `##$devianceResiduals`
 - `## Min Max`
 - `## -1.307112 1.412532`
 - `etc.....`

GLM model evaluation

- `predict()`
 - `> preds <- predict(model, training)`
 - `> errors <- select(`
 - `preds, preds$label, preds$prediction, preds$aircraft_type,`
 - `alias(preds$label - preds$prediction, "error"))`

GLM model vizualization

- `matplotlib`
- `Ggplot2`
- Third party libraries such as `prettyplotlib`